

# RISC-V and Cryptography

Developing XCrypto: An Embedded Class Crypto ISE

**Ben Marshall**, Daniel Page, Thinh Pham

University of Bristol  
Computer Science Department

# Outline

- ✿ Context and terms
- ✿ Implementing Cryptography on RISC-V
- ✿ XCrypto Overview
- ✿ Overlaps between XCrypto, P and B Extensions
- ✿ Side channels
- ✿ Cryptography on RISC-V: Where next?

---

## Context

- ✿ We are funded by the UK government ([EPSRC Grant EP/R012288/1](#)).
- ✿ All of our work is open source / free to use / MIT licensed / [on Github](#).
- ✿ An overarching project goal is to develop a side-channel resilient RISC-V CPU core.
- ✿ Along the way, we are doing lots of research into what decisions you can make at a (micro-)architectural level that help/hinder side channel resilience.
- ✿ We have focused on “embedded” class / IoT cores.
- ✿ Concretely, we are concerned with RV32 as a base ISA.

# Terms

- ✿ **Side channel:** measurements of a device / algorithms runtime, power consumption or electromagnetic energy emissions used to undermine a cryptographic construction.
- ✿ **Embedded CPU core:** Think ARM Cortex-M class, or anything with an IoT sticker. Probably doesn't have a cache. *Probably doesn't do macro-op fusion.*

# Cryptography on RISC-V

## The Bad News:

- ✿ The base ISA has a way to go to just be *on par* with peer ISAs.

## The Good News:

- ✿ We know exactly what functionality is needed to put it back on par.
- ✿ We know what is needed to make it *better than the competition*.
  - ▶ Lots of this has been independently developed in other extensions.
  - ▶ XCrypto has some helpful functionality in it not present in other extensions.
- ✿ There is **a lot** of prior art in this area.
  - ▶ Academic conferences like CHES feature existing research which has been under-used or re-invented.

---

# Implementing Crypto on RISC-V

- ✿ Considering the “general” or otherwise common RISC-V architecture strings:
  - ▶ Application Class: RV32G, RV64G (possibly with C too.)
  - ▶ Microcontrollers: RV32IMC
  
- ✿ These base architectures are deficient compared to ARM, MIPS in terms of performance, static and dynamic code density for most varieties of cryptographic workload.
  
- ✿ RISC-V puts the **R** in RISC.
  1. Lack of a rotate instruction.
  2. Limited memory addressing modes.
  3. Long arithmetic / integer carry support.
  4. Rigid adherence to 2-read-1-write.

# 1. Lack of a rotate instruction

Rotate is extremely popular in block ciphers and hashing algorithms.

ARX (Add, *Rotate*, XOR) constructions are common in light weight block ciphers designed for IoT devices.

By default, rotate is a 3-instruction sequence on RISC-V.

It's hard to overstate the impact of this:

Algorithm	Cycle Count		Gain
	Without Rotate	With Bitmanip Rotate	
SHA256	4205	3077	1.37x
ChaCha20	4149	2381	1.75x

As measured using Spike on RV64G, with(out) the draft B extension.

This will affect all RISC-V UNIX class systems going forward.

---

## 2. Limited addressing modes

- ✦ RISC-V does not support indexed addressing.
- ✦ This leads to *lots* of pointer arithmetic when stepping through arrays.
- ✦ Block ciphers/hash functions need to step non-linearly through state: AES/SHA3 are good examples.
  - ▶ In performance optimised crypto workloads for larger systems, you can unroll the loops and the problem goes away.
  - ▶ In embedded class cores, where loop unrolling is not an option, you see inner loops dominated by pointer arithmetic.
- ✦ Long-arithmetic also becomes dominated by pointer arithmetic.
- ✦ An argument against indexed load/store: the store variant requires three register read ports.
  - ▶ Counter argument: even the ARM M0+ has these sorts of instructions.



---

## 3. Long arithmetic support

- ✿ No carry detection in RISC-V, requires extra `slltu` instruction.
  - ▶ Lack of carry flags is nice sometimes, but this is the cost.
- ✿ ARM has `umlal` instructions: unsigned long multiply with accumulate. Present on ARM Cortex-M0 and up.
  - ▶ No equivalent on RISC-V
- ✿ Long arithmetic typically steps through three arrays at once: two sources and a destination.
  - ▶ Lack of indexed load means two extra instructions per inner loop iteration to increment pointers.

## 4. Rigid adherence to 2-read-1-write

- ✦ Having very simple encoding schemes is great for lots of reasons.
  - ▶ Lots of instructions only need to read two operands and write one result.
  - ▶ Fewer ports on a register file is nice, *but not as nice as it was 20+ years ago*.
- ✦ Lots of cryptographic functions are fundamentally reduction operations:
  - ▶ A plaintext *and* a key, reduced to a ciphertext.
  - ▶  $N$  bytes hashed down to  $M < N$  bytes.
- ✦ The more data you can feed to instructions, the faster reductions will run.
- ✦ At the ISA level: add more operands to your instructions.
- ✦ R.B. Lee, X. Yang, and Z. Shi. Validating Word-Oriented Processors for Bit and Multi-word Operations. In: Annual Computer Security Applications Conference (ACSAC). 2004, pp. 473488

---

## Implementing Crypto on RISC-V: Summary

- ✦ Cryptographic workloads are an excellent way to identify the shortcomings of RISC-V when compared to similar architectures.
- ✦ These shortcomings are not always cryptography specific, though cryptography is disproportionately affected.

---

# XCrypto: our custom crypto extension to RISC-V

XCrypto is designed to:

- ✦ Address the shortcomings in RISC-V already described.
- ✦ See *what else* we could add to RISC-V to make it as efficient at executing cryptographic workloads as possible.
- ✦ Be used as a vehicle to research what (*micro-*)*architectural* decisions you can make to help/hinder side channel resilience.
- ✦ Enable software based side channel countermeasures in a tighter performance budget.

XCrypto is not:

- ✦ A “drop in” embedded class crypto extension for RISC-V. (But we think it might guide the design of one...)

---

# XCrypto: Key design choices and assumptions

- ✦ Doesn't use the V extension as a base.
- ✦ Tries to add RISC instructions with general utility over very efficient but specialised CISC instructions.
- ✦ Adds an extra 16-element, 32-bit wide register file.
  - ▶ Assumes 3 read ports and 2 write ports.
  - ▶ Additional instructions do compute on this register file.
  - ▶ Memory address arithmetic still done in the normal GPRs.
- ✦ Tries to maximise implementation options:
  - ▶ Multi-cycle v.s. pipelined.
  - ▶ Integrated data path v.s Separate *secure* datapath.
  - ▶ In turn, maximises opportunities to implement side-channel countermeasures while minimising their cost.

---

# XCrypto: New Instructions

XCrypto includes instructions to address all of the problems we identified earlier:

✶ Rotate!

✶ Indexed and immediate offset Load/Store Unsigned Byte/Halfword/Word

- ▶ All XCrypto load/store instructions have *update semantics*. This lets us load a register with bytes/halfwords directly without shift-and-or sequences.

✶ Multi-operand instructions for long-integer arithmetic.

- ▶ Add/subtract, Multiply, Accumulate, Compare

---

# XCrypto: New Instructions

## General Purpose Instructions:

- 🔥 **Randomness:** Seeding, sampling and quality checking of entropy source.
- 🔥 **Memory:** Sub-word scatter/gather operations.
- 🔥 **Bit:** permutations, bitfield insert/extract, lookup-tables
- 🔥 **Packed:** SIMD-within-a-register (SWAR) arithmetic operations on 32/16/8/4/2 bit packed values.

## Special Purpose Instructions:

- 🔥 **AES:** lightweight AES sub-bytes and mix-columns acceleration
- 🔥 **SHA3:** code-dense, energy efficient state index generation

---

## XCrypto: New Instructions - Randomness

- ✿ A good source of randomness is essential for cryptography.
- ✿ Many side-channel countermeasures assume fast access to randomness.
- ✿ We don't specify how the randomness is generated, only that it is *cryptographically appropriate*.
- ✿ We only define an instruction based *interface* to a randomness source. That might be implemented as an alias for a CSR read, or memory mapped peripheral access.
- ✿ Instructions can seed a generator, sample the generator, and check if the generator thinks it currently has enough entropy to be useful.



## XCrypto: New Instructions - Scatter/Gather

- These instructions split a register into bytes or halfwords, and use each byte/halfword as an offset into memory from a base address.
- Useful for code-dense S-Box implementation.
- Amenable to hiding based side-channel countermeasures if defined appropriately.
- Still allows for masked S-Boxes.

`xc.gather.b crd, rs1, crs2:`

**begin**

$addr_0 \leftarrow GPR[rs1] + XCR[crs2]_0^8$

$addr_1 \leftarrow GPR[rs1] + XCR[crs2]_1^8$

$addr_2 \leftarrow GPR[rs1] + XCR[crs2]_2^8$

$addr_3 \leftarrow GPR[rs1] + XCR[crs2]_3^8$

$XCR[crd]_0^8 \leftarrow MEM[addr_0]$

$XCR[crd]_1^8 \leftarrow MEM[addr_1]$

$XCR[crd]_2^8 \leftarrow MEM[addr_2]$

$XCR[crd]_3^8 \leftarrow MEM[addr_3]$

**end**

---

## XCrypto: New Instructions - Bitwise

- ✿ Adds instructions for bit permutation, field insert/extract. All of which has been independently developed/added in the bitmanip extension.
- ✿ We also add support for very efficient 4-bit S-Box/LUTs and generalised ternary bitwise expressions.
- ✿ `xc.bop` (bit op) can implement any 3-variable bitwise expression. The example below is taken from SHA2.
- ✿ `xc.bop c0, c1, c2, 0xE8 // c0 = (c1&c2) | ((c1|c2) & c0)`

---

# XCrypto: New Instructions - Packed Arithmetic

- ✿ We add the standard set of arithmetic operations:
  - ▶ Add/sub/multiply/shift/rotate
  - ▶ Plus carryless multiply.
- ✿ All of these work as a SIMD instruction, on data elements 32, 16, 8, 4, or 2 bits wide.
- ✿ This is similar / identical to what the (DS)P extension proposes.
- ✿ We support narrower bit-widths to accelerate bit-sliced or secret shared implementations, designed for side-channel resilience.

---

# XCrypto: New Instructions - AES

- ✦ Light weight acceleration of the sub-bytes and mix-columns steps of AES.
- ✦ Based on existing academic literature:
  - ▶ Tillich, Stefan, and Johann Groschdl. "Instruction set extensions for efficient AES implementation on 32-bit processors." International workshop on cryptographic hardware and embedded systems. Springer, Berlin, Heidelberg, 2006.

---


## XCrypto: New Instructions - SHA3

XCrypto does a good job with SHA2. What about SHA3?

- ✂ The compute operations in SHA3 are simple: rotations and XORs.
- ✂ An unrolled implementation is 4K of code on RV32IMC.
- ✂ Unrolling is not realistic in embedded environments.
- ✂ Rolled-up implementations dominated by generating state matrix indexes:
  - ▶  $\text{index}(x, y) = (x\%5) + 5 * (y\%5)$
- ✂ Add novel instructions to generate the indexes from loop counter variables.
- ✂ Results in fast, code dense implementation with no extra lookup tables.
- ✂ Best performance increase per logic gate added we could find.
- ✂ 2x faster and 40% smaller across parameter sets.

# XCrypto: Runtime Improvement


XCrypto Runtime Improvement (Higher = Better)



As measured using the SPIKE ISA simulator and our area optimised reference implementation. We use hand-optimised RISC-V assembly as a baseline, unless the compiler is better than the human.

# XCrypto: Static Code Size

XCrypto Static Code Size Improvement (Lower = Better)



---

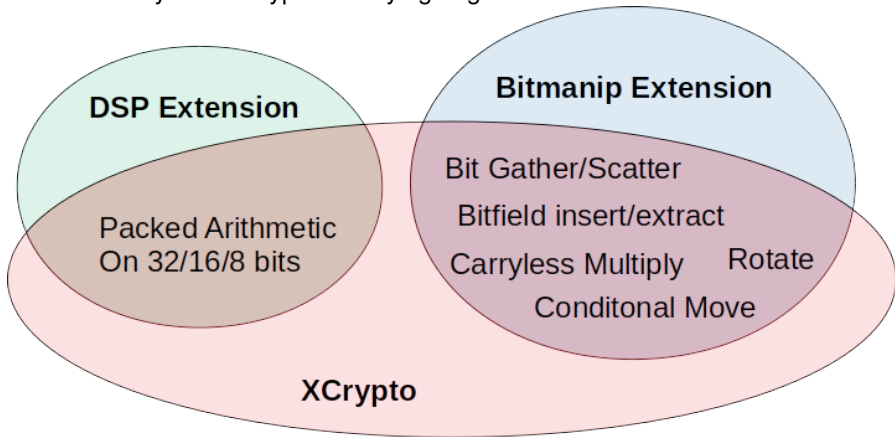
# XCrypto: Implementation Overhead

- ✶ We have a reference area-optimised Verilog implementation of XCrypto, which uses PicoRV32 as a base. When implemented on an FPGA:
  - ▶ XCrypto itself is 1.9K LUTs / 215 FFs Plus some DMEM slices.
  - ▶ The entire PicoRV32+XCrypto subsystem is 6.4K LUTs and 1.1K DFFs.
  - ▶ This is with no datapath sharing, and a DSP slices for the multiplier.
  - ▶ Most of the datapaths could be shared.
  - ▶ In an abstract CMOS gate flow, the area is 50% single-cycle multiplier.
- ✶ Takeaway: you can make this really small.



## Overlapping Functionality:

We've found that the DSP extension, and the Bitmanip extension overlap in terms of functionality with XCrypto to varying degrees.



## Side Channels

Side channel attacks are increasingly important:

- ✿ IoT implies lots of small devices needing to communicate.
- ✿ Communication implies a need for cryptography.
- ✿ IoT-class devices are very susceptible to side-channel attacks.

**Question:** Which decisions I make at an architectural level affect side channel resilience and how?

- ✿ Overwrite Semantics of Instructions
- ✿ Instruction Specification
- ✿ Implementation Flexibility
- ✿ Software Countermeasure Acceleration

---

# The RISC-V Crypto Extension: Our View

- ✿ It makes perfect sense to have whole/all-round SHA2/AES instructions *iff*:
  - ▶ Your implementation has the vector extension already.
  - ▶ Your implementation has enough vector lanes.
- ✿ There are a lot of cores which sit below this performance point *which need to do crypto efficiently*.
  - ▶ This class of core is the kind most likely to be subject to power and EM side-channel attacks.
  - ▶ They also most need to consider energy efficiency.
  - ▶ Currently, the public proposals regarding the crypto extension do not solve this problem.
- ✿ A RISC-V crypto extension should cater to all classes of core.

# Cryptography on RISC-V - What now?

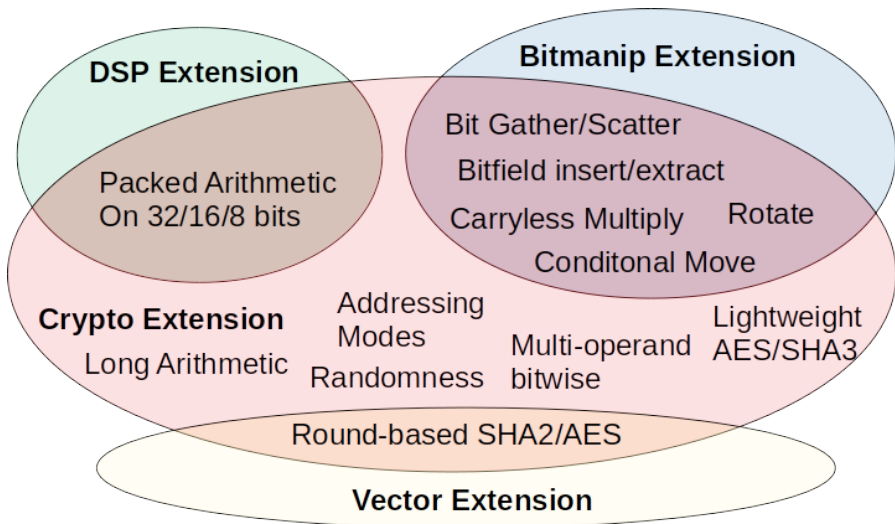
## Current Trajectory:

- ✿ Large Cores: Base ISA + Bitmanip + Vector Crypto
- ✿ Small Cores:
  - ▶ Base ISA + Bitmanip
  - ▶ Base ISA + *Custom Extensions*

## Possible Alternative:

- ✿ Large Cores: Base ISA + Bitmanip + Vector Crypto
- ✿ Small Cores: Base ISA + *Embedded Profile Crypto Extension*
- ✿ All Cores: Base ISA + *Modular Crypto Extension*

# Modular Crypto ISE



---

# Embedded Profile Crypto ISE

XCrypto was designed to be an embedded profile Crypto ISE.

Going forward, there are some possible options:

- ✦ Leave it as a custom extension:
  - ▶ It's a drop in, but with no compiler support etc.
- ✦ Do some work to turn it into a standard extension:
  - ▶ Encoding Space, Extra State Choices etc.
  - ▶ Handling overlapping functionality.
  - ▶ Needs to fix a remit.
  - ▶ This can be done if there is an appetite for it?

# Cryptography on RISC-V

## The Bad News:

- ✿ The base ISA has a way to go to just be *on par* with peer ISAs.

## The Good News:

- ✿ We know exactly what functionality is needed to put it back on par.
- ✿ We know what is needed to make it *better than the competition*.
  - ▶ Lots of this has been independently developed in other extensions.
  - ▶ XCrypto has some helpful functionality in it not present in other extensions.
- ✿ There is **a lot** of prior art in this area.
  - ▶ Academic conferences like CHES feature existing research which has been under-used or re-invented.
- ✿ **We are keen to help take all of this forward.**