

Securing SoCs from Time Side Channels


Jose Renau

Esperanto _____
Technologies _____

An Attacker can use Time Leaks to Infer Usage



IPC = 0.9
Time...
Misses...



IPC = 1.1
Time...
Misses...



IPC = 0.6
Time...
Misses...


Anatomy of an Attack

- Victim **runs code that leaks** observable side-effects.
 - Attacker runs code that is **affected by timing leaks**.
 - Attacker **measures time** his code took to run.
-
- Can we prevent to run code? (not clear)
 - Can we avoid time leaks? (the goal of this)
 - Can we avoid measuring time? (not really)


Time Domain

One Time Domain should not affect the performance of another Time Domain

Time Domain 1
IPC1




```
add  s0,sp,48
sw   a0,-36(s0)
sd   a1,-48(s0)
lw   a5,-36(s0)
sw   a5,-20(s0)
sw   zero,-24(s0)
j     .L2
```



Time Impact

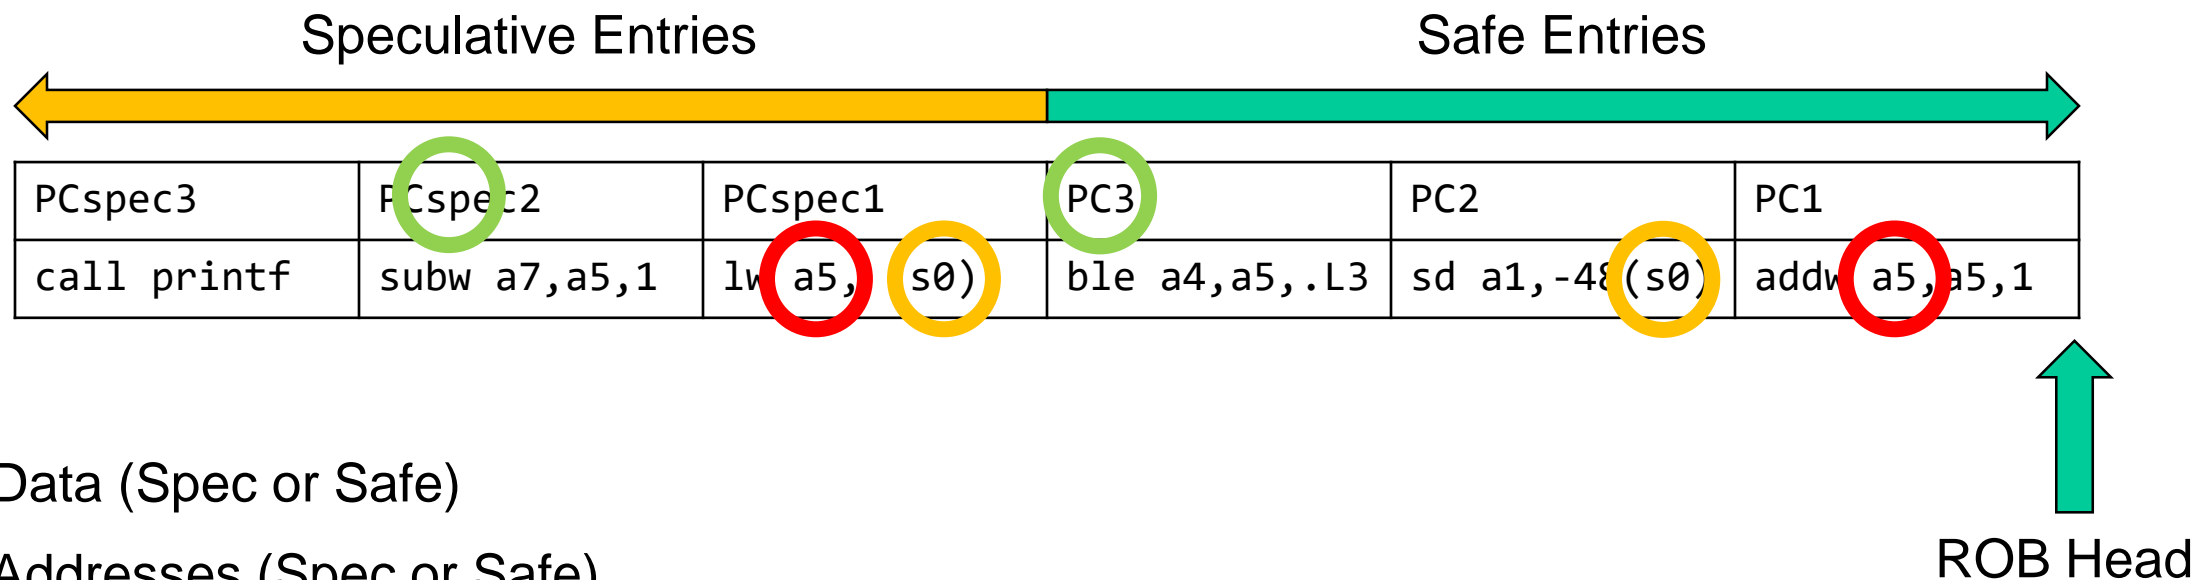
Time Domain 2
IPC2



```
lw   a4,-20(s0)
lw   a5,-36(s0)
addw a5,a4,a5
sw   a5,-20(s0)
lw   a5,-24(s0)
addw a5,a5,1
sw   a5,-24(s0)
```

- Categorize Time Leaks
- High level techniques
- Sample cores

What can Time Leak?



D: Data (Spec or Safe)

A: Addresses (Spec or Safe)

P: Program Counter (Spec or Safe)

E: Execution Time (Spec or Safe)

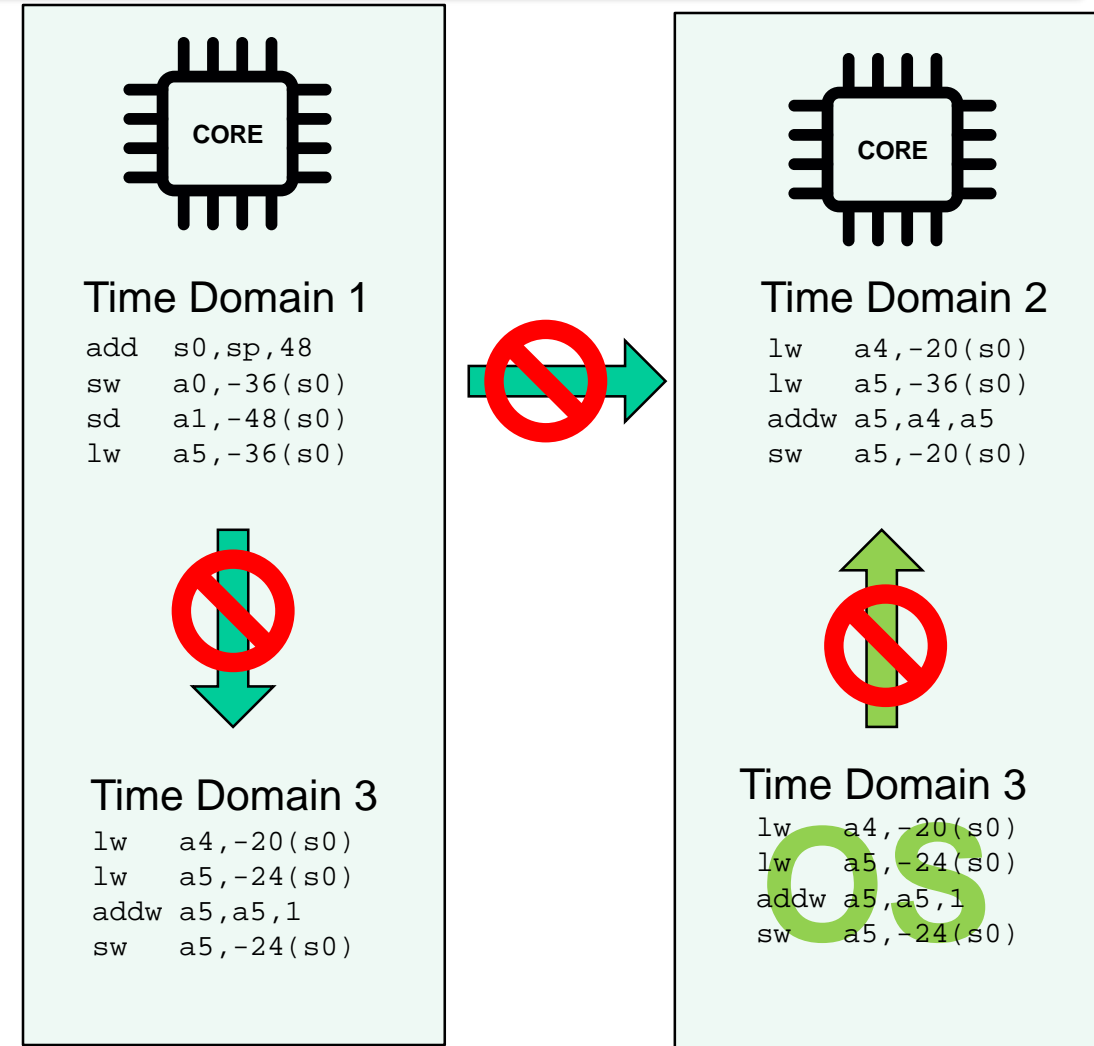
Where can Time Leaks go?

SC: no Same Core leaks

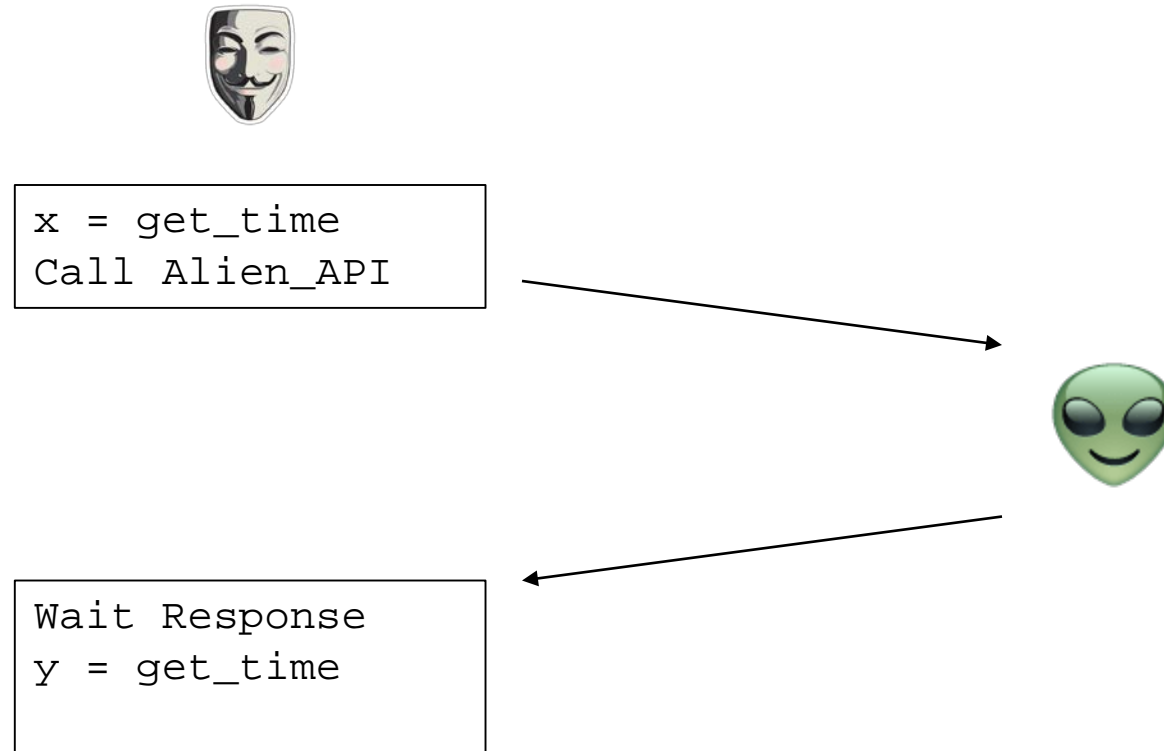
MC: no Multi-Core leaks

OS: no Operating System leaks

AI: no Application Interface leaks



What is a Safe.AI? (Application Interface)



Time Leaks Categorization

		Where Leaks			
		SC	MC	OS	AI
What Leaks	D	✓	✓	✓	✓
	A				
	P				
	E				

Safe[D].ALL core

		Where Leaks			
		SC	MC	OS	AI
What Leaks	D	✓			
	A	✓			
	P	✓			
	E	✓			

Safe.SC core

Time Leaks Categories

		Where Leaks			
		SC	MC	OS	AI
What Leaks	D		✓	✓	
	A		✓	✓	
	P		✓	✓	
	E		✓	✓	

Safe.MC.OS

		Where Leaks			
		SC	MC	OS	AI
What Leaks	D	✓	✓	✓	✓
	A	✓	✓	✓	✓
	P	✓	✓	✓	✓
	E	✓	✓	✓	✓

Spec.ALL

Interesting Cases

- Google/amazon may want
 - Safe.MC.OS, Spec.ALL
- Desktop/Phone
 - Safe.SC.MC, Spec.ALL
- Autonomous cars
 - Spec.MC, Spec.ALL and use OS to separate per core
- Router/switch
 - Safe.nothing?
- IOT
 - Unclear, web browser IOT, maybe Spec.??, Spec.ALL

- Categorize Time Leaks
- High level techniques
- Sample cores

High Level Idea

- For Spec:
 - Do not make speculation visible
- For Safe:
 - Do not make any traffic visible to other time domains

Tech 1: Point of No Return (PNR)

- Not all the instructions are speculative
- PNR in ROB indicates point where instruction is guaranteed to eventually commit
- SPEC2006 for an A72-like core ~25% of the instructions are beyond the PNR

Tech 2: No Speculative Update Predictors

- Delay and/or perfect fixup all the predictors
 - Good: Minor/no performance impact
 - Bad: More buffering
- Delay TAGE/BTB/... updates until retirement
 - No performance impact!
- Return Address Stack (RAS) and GHR
 - Have a copy at retirement, perfect fix at missprediction
- Delay AMPM/Stride/xxx prefetch updates to retirement
 - Even better performance, no pollution

Where Leaks				
	SC	MC	OS	AI
D	✓	✓	✓	✓
A	✓	✓	✓	✓
P	✓	✓	✓	✓
E				

SPEC

Tech 3: No Speculative Cache

- No inclusive/exclusive L1/L2/L3 caches
- Allocate lines only to Store Completion Buffer (SCB) on cache miss
 - When combined with PNR, very few entries in SCB
 - For a 4-way OoO core and SPEC2006, 3 entries cover over 99% of the cases
- Issues: Request does not change any state until passes PNR
 - If it triggers coherence, it has to wait until passes PNR
 - Cache LRU updated at retirement, not execute
- Possible to add “speculative buffers” for each cache

Where Leaks

	SC	MC	OS	AI
D	✓	✓	✓	✓
A	✓	✓	✓	✓
P	✓	✓	✓	✓
E				

SPEC

Tech 4: Constant Miss-Speculation Recovery

- No resources pin down by speculative flushes
 - No long latency resource pindown (sqrt/div/...)
 - No miss queue/MSHR

Where Leaks

	SC	MC	OS	AI
D	✓	✓	✓	✓
A	✓	✓	✓	✓
P	✓	✓	✓	✓
E	✓	✓	✓	✓

SPEC

Tech 5: Partition and BW isolation

- Private L1 and L2
- Many papers showing how to data partition LLC
- Bias designs towards systems with smaller amount of shared LLC
- Partition the directory, LLC, and any resource

Where Leaks

	SC	MC	OS	AI
D	✓	✓	✓	✓
A	✓	✓	✓	✓
P	✓	✓	✓	✓
E	✓	✓	✓	✓

SPEC

Tech 6: Partial/Full Flush

- Flush entries at Time Domain switch
- Flush can be partial like 2 LRU ways or X entries

Tech 7: Partial/Full Off-load

- Off-load/Re-load part of structures at Time Domain switch
- E.g:
 - Off-Load Data L1 MRU lines on Time Domain switch
 - Flush the other Data L1 entries

Tech 8: Dynamic every ?M cycles

- Dynamic partition/repartition/adaptation can increase throughput/utilization
- OK to dynamically adapt as long as it does not leak
 - Adapt to a slow moving average infrequently (millions of cycles)
 - Trade-off performance vs security

Other Techniques (not show due to time)

- Tech 9: Dealing with Row Hammer
- Tech 10: BW isolation per Time Domain
- Tech 11: QoS for network on chip
- Tech 12: Non-value dependent operation latency (no telescopic units)
- Tech 13: Regularization
- Tech 14: Homogenization

- Categorize Time Leaks
- High level techniques
- Sample cores

Spec.ALL Techniques

- Overall different structure
 - Private L1/L2 non-inclusive caches, small non-inclusive L3
 - Extra “fake” physical address bits to indicate Time Domain id (core and OS/User)
 - Inclusive L1 and L2 TLB. Fast L2 to L1 TLB update
 - L2 instruction cache (IL2)
- Core
 - Tech 1: PNR
 - Tech 2: No Speculative Update Predictors
 - TAGE/ITTAGE/BTB/Prefetcher, stride prefetcher uses retire address + extra offset
 - Tech 3: No Speculative Cache
 - Non inclusive L1/L2/L3. Allocate on L1 SCB
 - Tech 4: Constant spec flush time
- SoC
 - Tech 5: Dynamic Partition
 - Fixed partition LLC, directory
 - Dynamic partition: bandwidth, Memory Controller (open pages, row hammer)
 - Tech 10: Memory BW partitioning per core

Safe.MC Techniques

- Comes for free once you have the previous Spec.ALL
 - Tech 5 (Dynamic Partitioning) and Tech 10 (BW)

Mostly, Spec.ALL already provides Safe.MC

Safe.OS Techniques

- Tech 5: Dynamic Partition if frequent OS calls
 - Dynamic partition: L2 TLB, L2, Memory Dependence Predictor
 - No L2 allocate by OS unless frequent OS
- Tech 6: Full/Partial Flush on OS return
 - Memory Dependence predictor, L1 TLB, Memory Controller
 - Data L1, Instruction L1, BTB, ITTAGE
- Tech 8: Dynamic adapt every SoC partitioned resources every 10M cycles
 - L2, L2 TLB, BTB, ITTAGE

Mostly, allocate a chunk of core predictors to OS or flush a fixed chunk at OS return

Safe.SC.MC.OS Techniques

- Combine Safe.MC and Safe.OS techniques, but also add
- Tech 5: Dynamic Partition
 - Dynamic partition: L2 TLB, DL2, IL1
- Tech 6: Partial Flush if infrequent Time Domain switch
 - Full Flush Memory Dependence predictor, L1 TLB
 - Partial flush if partition DL2, IL2, BTB, ITTAGE
- Tech 7: Partial Off-load
 - Off-load MRU sets in data and instruction DL1 and IL1 on Time Domain switch
 - Off-load to IL1/DL2 if partition mode, flush otherwise
 - Off-load data and instruction prefetch tables for quick re-load

Mostly, flush if infrequent, partition/off-load if frequent switching

Conclusions

- Categorization of Time Leaks to classify/understand/test current SoCs
- Show several techniques that can achieve different degrees of protection
- Show techniques for Spec.ALL with small performance impact (but large redesign)
- Show techniques for Safe.OS, Safe.MC, Safe.SC
- Protection level could change at run-time

What do we need from ISA?


- Some way to indicate TimeDomain ID
- Option 1
 - Have CSR
- Option 2
 - Use the upper bits in the address space
- TimeDomain ID is important to avoid overheads when unnecessary
 - E.g: no need to wait for PNR in most data sharing multithreaded apps
- Some way to expose the reconfiguration/repartitioning. An AMPM++?

Questions?

Jose Renau

*Department of Computer Engineering,
University of California, Santa Cruz*

<http://masc.soe.ucsc.edu>



- [Securing Processors from Time Side Channels](#), Jose Renau, April 2018.